

Assignment:-

Applying GBDT and RF on Amazon Fine Food Reviews Analysis

Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews> (<https://www.kaggle.com/snap/amazon-fine-food-reviews>)

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

1. Objective:

Given a review, determine whether the review is positive Rating (4 or 5) or negative rating (1 or 2). Use BoW, TF-IDF, Avg-Word2Vec, TF-IDF-Word2Vec to vectorise the reviews. Apply GBDT and Random Forest Algorithm for Amazon fine food Reviews find right baselearners using cross validation Get feature importance for positive class and Negative class

In [1]:

```
# Loading required libraries
import warnings
warnings.filterwarnings('ignore')

import numpy as np
import pandas as pd
import matplotlib
import sqlite3
import string
import gensim
import scipy
import nltk
import time
import seaborn as sns
from scipy import stats
from matplotlib import pyplot as plt

from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer

from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from sklearn import metrics
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_curve, roc_auc_score, auc
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_recall_fscore_support as prf1

from sklearn.model_selection import KFold
from sklearn.model_selection import train_test_split
```

1.1 Connecting SQL file

In [2]:

```
#Loading the data
con = sqlite3.connect('./final.sqlite')

data = pd.read_sql_query("""
SELECT *
FROM Reviews
""", con)
```

In [3]:

```
print(data.shape)
data.head()
```

(364171, 12)

Out[3]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Helpful
0	138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski	0	
1	138688	150506	0006641040	A2IW4PEEKO2R0U	Tracy	1	
2	138689	150507	0006641040	A1S4A3IQ2MU7V4	sally sue "sally sue"	1	
3	138690	150508	0006641040	AZGXZ2UUK6X	Catherine Hallberg " (Kate)"	1	
4	138691	150509	0006641040	A3CMRKGE0P909G	Teresa	3	



1.2 Data Preprocessing

In [4]:

```
data.Score.value_counts()  
#i had done data preprocessing i had stored in final.sqlite now loaded this file no need to
```

Out[4]:

positive 307061
negative 57110
Name: Score, dtype: int64

1.3 Sorting the data

In [5]:

```
# Sorting the data according to the time-stamp  
sorted_data = data.sort_values('Time', axis=0, ascending=True, inplace=False, kind='quicksort')  
sorted_data.head()
```

Out[5]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpDenominator
0	138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski	0	1
30	138683	150501	0006641040	AJ46FKXOVC7NR	Nicholas A Mesiano	2	3
424	417839	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	0	1
330	346055	374359	B00004CI84	A344SMIA5JECGM	Vincent P. Ross	1	2
423	417838	451855	B00004CXX9	AJH6LUC1UT1ON	The Phantom of the Opera	0	1

1.4 Mapping

In [6]:

```
def partition(x):
    if x == 'positive':
        return 1
    return 0

#Preparing the filtered data
actualScore = sorted_data['Score']
positiveNegative = actualScore.map(partition)
sorted_data['Score'] = positiveNegative
sorted_data.head()
```

Out[6]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Help
0	138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski	0	
30	138683	150501	0006641040	AJ46FKXOVC7NR	Nicholas A Mesiano	2	
424	417839	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	0	
330	346055	374359	B00004CI84	A344SMIA5JECGM	Vincent P. Ross	1	
423	417838	451855	B00004CXX9	AJH6LUC1UT1ON	The Phantom of the Opera	0	

1.5 Taking First 150k rows

In [7]:

```
# We will collect different 150000 rows without repetition from time_sorted_data dataframe
my_final = sorted_data[:150000]
print(my_final.shape)
my_final.head()
```

(150000, 12)

Out[7]:

	index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Help
0	138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski	0	
30	138683	150501	0006641040	AJ46FKXOVC7NR	Nicholas A Mesiano	2	
424	417839	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	0	
330	346055	374359	B00004CI84	A344SMIA5JECGM	Vincent P. Ross	1	
423	417838	451855	B00004CXX9	AJH6LUC1UT1ON	The Phantom of the Opera	0	



1.6 Splitting data into train and test based on time (70:30)

In [8]:

```

from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_validate

x=my_final['CleanedText'].values
y=my_final['Score']

#Splitting data into train test and cross validation
x_train,x_test,y_train,y_test =train_test_split(x,y,test_size =0.3,random_state = 42)

print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)

```

```

(105000,)
(45000,)
(105000,)
(45000,)

```

2. Techniques For Vectorization

Why we have to convert text to vector

By converting text to vector we can use whole power of linear algebra.we can find a plane to seperate

2.1 BOW

In [9]:

```

#Bow

from sklearn.feature_extraction.text import CountVectorizer
count_vect = CountVectorizer()
final_counts_Bow_tr= count_vect.fit_transform(x_train)# computing Bow
print("the type of count vectorizer ",type(final_counts_Bow_tr))
print("the shape of out text BOW vectorizer ",final_counts_Bow_tr.get_shape())
print("the number of unique words ", final_counts_Bow_tr.get_shape()[1])
final_counts_Bow_test= count_vect.transform(x_test)# computing Bow
print("the type of count vectorizer ",type(final_counts_Bow_test))
print("the shape of out text BOW vectorizer ",final_counts_Bow_test.get_shape())

```

```

the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (105000, 38300)
the number of unique words 38300
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (45000, 38300)

```

2.2 Normalizing Data

In [10]:

Data-preprocessing: Normalizing Data

```

from sklearn import preprocessing
standardized_data_train = preprocessing.normalize(final_counts_Bow_tr)
print(standardized_data_train.shape)
standardized_data_test = preprocessing.normalize(final_counts_Bow_test)
print(standardized_data_test.shape)

```

(105000, 38300)

(45000, 38300)

2.3 Applying RandomForest Algorithm

In [11]:

#Gridsearch Cross Validation

```

from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn import tree

base_learners = [40,60,80,100,120]
depth=[3,5,7,9,11]
param_grid = {'n_estimators': base_learners, 'max_depth':depth}
rf = RandomForestClassifier(min_samples_leaf=5,max_features='sqrt',criterion='gini',random_
model = GridSearchCV(rf, param_grid,scoring = 'f1',cv=3 , n_jobs = -1,pre_dispatch=2)
model.fit(standardized_data_train,y_train)
print(model.best_score_, model.best_params_)
print("Model with best parameters : \n",model.best_estimator_)
print("Accuracy of the model : ",model.score(standardized_data_test, y_test))

a = model.best_params_
optimal_estimator = a.get('n_estimators')
optimal_depth = a.get('max_depth')

```

0.9169864843830612 {'max_depth': 11, 'n_estimators': 120}

Model with best parameters :

```

RandomForestClassifier(bootstrap=True, class_weight='balanced',
                        criterion='gini', max_depth=11, max_features='sqrt',
                        max_leaf_nodes=None, min_impurity_decrease=0.0,
                        min_impurity_split=None, min_samples_leaf=5,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=120, n_jobs=None, oob_score=False,
                        random_state=100, verbose=0, warm_start=False)

```

Accuracy of the model : 0.9161176281536236

In [12]:

```


results = model.cv_results_
meanscore=results['mean_test_score']

```

Heatmap for plotting CV Scores

In [13]:

```
pvt = pd.pivot_table(pd.DataFrame(model.cv_results_), values='mean_test_score', index='param_
import seaborn as sns
ax = sns.heatmap(pvt, annot=True, fmt="f")
```



In [14]:


```
clf = RandomForestClassifier(n_estimators=optimal_estimator, class_weight='balanced')
clf.fit(standardized_data_train, y_train)
y_pred = clf.predict(standardized_data_test)
```

2.4 Confusion Matrix

In [15]:

```
cm_bow=confusion_matrix(y_test,y_pred)
print("Confusion Matrix:")
sns.heatmap(cm_bow, annot=True, fmt='d')
plt.show()
```

Confusion Matrix:



In [16]:

```
#finding out true negative , false positive , false negative and true positive
tn, fp, fn, tp = cm_bow.ravel()
( tp, fp, fn, tp)
print(" true negitves are {} \n false positives are {} \n false negatives are {} \n true pos
```

```
true negitves are 1061
false positives are 5053
false negatives are 56
true positives are 38830
```

2.5 Calculating Accuracy, Error on test data, Precision, Recall, Classification Report

In [17]:

```

from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import classification_report
from sklearn.metrics import roc_auc_score

# evaluating accuracy
acc_bow = accuracy_score(y_test, y_pred) * 100
print('\nTest Accuracy of the RandomForest for base_learners ={}, depth ={} is accuracy {}'.format(base_learners, depth, acc_bow))

# Error on test data
test_error_bow = 100-acc_bow
print("\nTest Error RandomForest for base learner and optimal_depth is %f%%" % (test_error_bow))

# evaluating precision
precision_score = precision_score(y_test, y_pred)
print('\nThe Test Precision of the RandomForest for base_learners ={}, depth ={} is accuracy {}'.format(base_learners, depth, precision_score))

# evaluating recall
recall_score = recall_score(y_test, y_pred)
print('\nThe Test Recall of the RandomForest for base_learners ={}, depth ={} is accuracy {}'.format(base_learners, depth, recall_score))

# evaluating Classification Report
classification_report = classification_report(y_test, y_pred)
print('\nThe Test classification report for base learner and optimal_depth is: \n\n', (classification_report))

```

Test Accuracy of the RandomForest for base_learners =120, depth =11 is accuracy 88.64666666666666

Test Error RandomForest for base learner and optimal_depth is 11.353333%

The Test Precision of the RandomForest for base_learners =120, depth =11 is accuracy 0.8848529043137434

The Test Recall of the RandomForest for base_learners =120, depth =11 is accuracy 0.9985598930206244


The Test classification report for base learner and optimal_depth

	precision	recall	f1-score	support
0	0.95	0.17	0.29	6114
1	0.88	1.00	0.94	38886
micro avg	0.89	0.89	0.89	45000
macro avg	0.92	0.59	0.62	45000
weighted avg	0.89	0.89	0.85	45000

2.6 Plotting roc_auc curve

In [18]:

```
y_pred_proba = clf.predict_proba(standardized_data_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
plt.plot(fpr,tpr,label="bow, AUC="+str(auc))
plt.plot([0,1],[0,1],'r--')
plt.title('ROC curve: Random Forest')
plt.legend(loc='lower right')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



2.7 Top 25 words

In [19]:

```
words = count_vect.get_feature_names()
likelihood_df = pd.DataFrame(clf.feature_importances_.transpose(), columns=[ 'Score'], index=
top_25 = likelihood_df.sort_values(by='Score', ascending=False).iloc[:25]
top_25.reset_index(inplace=True)
top_words = top_25['index']
print(top_words)
```

```
0      great
1       love
2       best
3  disappoint
4     delici
5       good
6       tast
7    perfect
8      would
9    product
10      bad
11     find
12    excel
13     like
14    money
15  favorit
16     wast
17  thought
18     use
19    didnt
20    easi
21    make
22    nice
23   wonder
24   flavor
Name: index, dtype: object
```

In [20]:

```
from wordcloud import WordCloud

list_of_words_str = ' '.join(top_words)

wc = WordCloud(background_color="white", max_words=len(top_words),
                width=900, height=600, collocations=False)
wc.generate(list_of_words_str)
print ("\n\nWord Cloud for Important features")
plt.imshow(wc, interpolation='bilinear')
plt.axis("off")
plt.show()
```

Word Cloud for Important features



3.Applying GBDT Algorithm

In [21]:

```

from sklearn.ensemble import GradientBoostingClassifier

Learning_rate = [0.05,0.1,0.2,0.3]
depth=[3,5,7,9,11]
param_grid = {'max_depth':depth, 'learning_rate':Learning_rate}
gb = GradientBoostingClassifier(loss='deviance',max_features='sqrt',subsample=0.1,n_estimators=100)
model = GridSearchCV(gb, param_grid, scoring = 'f1', cv=3 , n_jobs = -1,pre_dispatch=2)
model.fit(standardized_data_train,y_train)
print("Model with best parameters :\n",model.best_estimator_)
print("Accuracy of the model : ",model.score(standardized_data_test, y_test))
a = model.best_params_
optimal_learningrate = a.get('learning_rate')
optimal_depth = a.get('max_depth')

```

Model with best parameters :

```

GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.1, loss='deviance', max_depth=11,
                           max_features='sqrt', max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=100,
                           n_iter_no_change=None, presort='auto', random_state=None,
                           subsample=0.1, tol=0.0001, validation_fraction=0.1,
                           verbose=0, warm_start=False)

```

Accuracy of the model : 0.9409884855581576

In [22]:

```

results = model.cv_results_
meanscore=results['mean_test_score']

```


Heatmap for plotting CV Scores

In [23]:

```

pvt =pd.pivot_table(pd.DataFrame(model.cv_results_),values='mean_test_score', index='param_
import seaborn as sns
ax = sns.heatmap(pvt,annot=True,fmt="f")

```



In [24]:


```
clf = GradientBoostingClassifier(max_depth=optimal_depth,learning_rate=optimal_learningrate)
clf.fit(standardized_data_train,y_train)
y_pred = clf.predict(standardized_data_test)
```

3.1 Confusion Matrix

In [25]:

```
## Confusion Matrix:
cm_bow=confusion_matrix(y_test,y_pred)
print("Confusion Matrix:")
sns.heatmap(cm_bow, annot=True, fmt='d')
plt.show()
```

Confusion Matrix:



In [26]:

```
#finding out true negative , false positive , false negative and true positive
tn, fp, fn, tp = cm_bow.ravel()
( tp, fp, fn, tp)
print(" true negitves are {} \n false positives are {} \n false negatives are {} \n true positives are {}")
```

```
true negitves are 1405
false positives are 4709
false negatives are 287
true positives are 38599
```

3.2 Calculating Accuracy,Error on test data,Precision,Recall,Classification Report

In [27]:

```

from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import classification_report
from sklearn.metrics import roc_auc_score

# evaluating accuracy
acc_bow = accuracy_score(y_test, y_pred) * 100
print('\nTest Accuracy of the GBDT for maxdepth={}, Learning_rate ={} is accuracy {}'.format(

# Error on test data
test_error_bow = 100-acc_bow
print("\nTest Error GBDT for maxdepth and Learning_rate is  %f%%" % (test_error_bow))

# evaluating precision
precision_score = precision_score(y_test, y_pred)
print('\nThe Test Precision of the GBDT for maxdepth={}, Learning_rate ={} is accuracy {}'.format(

# evaluating recall
recall_score = recall_score(y_test, y_pred)
print('\nThe Test Recall of the GBDT for maxdepth={}, Learning_rate ={} is accuracy {}'.format(

# evaluating Classification Report
classification_report = classification_report(y_test, y_pred)
print('\nThe Test classification report for GBDT maxdepth and Learning_rate  \n\n ',(class

```

Test Accuracy of the GBDT for maxdepth=11, Learning_rate =0.1 is accuracy 88.89777777777778

Test Error GBDT for maxdepth and Learning_rate is 11.102222%

The Test Precision of the GBDT for maxdepth=11, Learning_rate =0.1 is accuracy 0.8912672023644592

The Test Recall of the GBDT for maxdepth=11, Learning_rate =0.1 is accuracy 0.9926194517307


The Test classification report for GBDT maxdepth and Learning_rate

	precision	recall	f1-score	support
0	0.83	0.23	0.36	6114
1	0.89	0.99	0.94	38886
micro avg	0.89	0.89	0.89	45000
macro avg	0.86	0.61	0.65	45000
weighted avg	0.88	0.89	0.86	45000

3.3 Plotting roc_auc curve

In [28]:

```
y_pred_proba = clf.predict_proba(standardized_data_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
plt.plot(fpr,tpr,label="bow, AUC="+str(auc))
plt.plot([0,1],[0,1],'r--')
plt.title('ROC curve: GBDT')
plt.legend(loc='lower right')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



3.4 Top 25 words

In [29]:

```
words = count_vect.get_feature_names()
likelihood_df = pd.DataFrame(clf.feature_importances_.transpose(), columns=[ 'Score'], index=
top_25 = likelihood_df.sort_values(by='Score', ascending=False).iloc[:25]
top_25.reset_index(inplace=True)
top_words = top_25['index']
print(top_words)
```

```
0    disappoint
1         great
2        terribl
3         wast
4         bad
5        worst
6         love
7        horribl
8        thought
9         best
10         tast
11        would
12        threw
13        money
14         aw
15        didnt
16        return
17        delici
18        stale
19        receiv
20        wouldnt
21        gross
22         poor
23         list
24        product
Name: index, dtype: object
```

In [30]:

```

from wordcloud import WordCloud

list_of_words_str = ' '.join(top_words)

wc = WordCloud(background_color="white", max_words=len(top_words),
               width=900, height=600, collocations=False)
wc.generate(list_of_words_str)
print("\n\nWord Cloud for Important features")
plt.imshow(wc, interpolation='bilinear')
plt.axis("off")
plt.show()

```

Word Cloud for Important features



4. TF-IDF

In [31]:

```

#tf-idf
from sklearn.feature_extraction.text import TfidfVectorizer
tf_idf_vect = TfidfVectorizer()

final_counts_tfidf_tr= tf_idf_vect.fit_transform(x_train)
print("the type of count vectorizer ",type(final_counts_tfidf_tr))
print("the shape of out text tfidf vectorizer ",final_counts_tfidf_tr.get_shape())
print("the number of unique words ", final_counts_tfidf_tr.get_shape()[1])
final_counts_tfidf_test= tf_idf_vect.transform(x_test)
print("the type of count vectorizer ",type(final_counts_tfidf_test))
print("the shape of out text tfidf vectorizer ",final_counts_tfidf_test.get_shape())
print("the number of unique words ", final_counts_tfidf_test.get_shape()[1])

```

```

the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text tfidf vectorizer (105000, 38300)
the number of unique words 38300
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text tfidf vectorizer (45000, 38300)
the number of unique words 38300

```

4.1 Normalizing Data

In [32]:

```
# Data-preprocessing: Normalizing Data
from sklearn import preprocessing
standardized_data_train = preprocessing.normalize(final_counts_tfidf_tr)
print(standardized_data_train.shape)
standardized_data_test = preprocessing.normalize(final_counts_tfidf_test)
print(standardized_data_test.shape)
```

```
(105000, 38300)
(45000, 38300)
```

4.2 Applying RandomForest Algorithm

In [33]:

```
# Gridsearch Cross Validation

from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn import tree

base_learners = [40,60,80,100,120]
depth=[3,5,7,9,11]
param_grid = {'n_estimators': base_learners, 'max_depth':depth}
rf = RandomForestClassifier(min_samples_leaf=5,max_features='sqrt',criterion='gini',random_
model = GridSearchCV(rf, param_grid,scoring = 'f1',cv=3 , n_jobs = -1,pre_dispatch=2)
model.fit(standardized_data_train,y_train)
print(model.best_score_, model.best_params_)
print("Model with best parameters :\n",model.best_estimator_)
print("Accuracy of the model : ",model.score(standardized_data_test, y_test))

a = model.best_params_
optimal_estimator = a.get('n_estimators')
optimal_depth = a.get('max_depth')
```

```
0.9127876758692683 {'max_depth': 11, 'n_estimators': 120}
Model with best parameters :
RandomForestClassifier(bootstrap=True, class_weight='balanced',
criterion='gini', max_depth=11, max_features='sqrt',
max_leaf_nodes=None, min_impurity_decrease=0.0,
min_impurity_split=None, min_samples_leaf=5,
min_samples_split=2, min_weight_fraction_leaf=0.0,
n_estimators=120, n_jobs=None, oob_score=False,
random_state=100, verbose=0, warm_start=False)
Accuracy of the model : 0.910332405876581
```

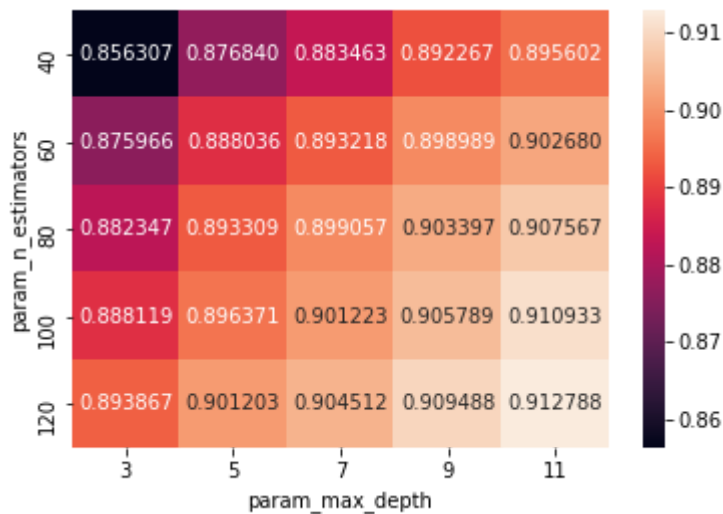
In [34]:

```
results = model.cv_results_
meanscore=results['mean_test_score']
```

Heatmap for Plotting CV Scores

In [35]:

```
pvt = pd.pivot_table(pd.DataFrame(model.cv_results_), values='mean_test_score', index='param_
import seaborn as sns
ax = sns.heatmap(pvt, annot=True, fmt="f")
```



In [36]:


```
clf = RandomForestClassifier(n_estimators=optimal_estimator, class_weight='balanced')
clf.fit(standardized_data_train, y_train)
y_pred = clf.predict(standardized_data_test)
```

4.3 Confusion Matrix

In [37]:

```
cm_tfidf = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
sns.heatmap(cm_bow, annot=True, fmt='d')
plt.show()
```

Confusion Matrix:



In [38]:

```
#finding out true negative , false positive , false negative and true positive
tn, fp, fn, tp = cm_tfidf.ravel()
( tp, fp, fn, tp)
print(" true negitves are {} \n false positives are {} \n false negatives are {} \n true pos
```

```
true negitves are 1139
false positives are 4975
false negatives are 67
true positives are 38819
```

4.4 Calculating Accuracy, Error on test data, Precision, Recall, Classification Report

In [39]:

```

from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import classification_report

# evaluating accuracy
acc_tfidf = accuracy_score(y_test, y_pred) * 100
print('\nThe Test Accuracy of the RandomForest for base_learners ={}, depth ={} is accuracy

# Error on test data
test_error_tfidf = 100-acc_tfidf
print("\nTest Error of the RandomForest for optimal_depth  %f%%" % (test_error_tfidf))

# evaluating precision
precision_score = precision_score(y_test, y_pred)
print('\nThe Test Precision of the RandomForest for base_learners ={}, depth ={} is accuracy

# evaluating recall
recall_score = recall_score(y_test, y_pred)
print('\nThe Test Recall of the RandomForest for base_learners ={}, depth ={} is accuracy {

# evaluating Classification report
classification_report = classification_report(y_test, y_pred)
print('\nThe Test classification report of the RandomForest for optimal_depth is \n\n ',(cl

```

The Test Accuracy of the RandomForest for base_learners =120, depth =11 is a ccuracy 88.79555555555557

Test Error of the RandomForest for optimal_depth 11.204444%

The Test Precision of the RandomForest for base_learners =120, depth =11 is accuracy 0.8863999634653149

The Test Recall of the RandomForest for base_learners =120, depth =11 is accuracy 0.9982770148639614


The Test classification report of the RandomForest for optimal_depth is

	precision	recall	f1-score	support
0	0.94	0.19	0.31	6114
1	0.89	1.00	0.94	38886
micro avg	0.89	0.89	0.89	45000
macro avg	0.92	0.59	0.63	45000
weighted avg	0.89	0.89	0.85	45000

4.5 Plotting roc_auc curve

In [40]:

```
y_pred_proba = clf.predict_proba(standardized_data_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
plt.plot(fpr,tpr,label="tfidf, AUC="+str(auc))
plt.plot([0,1],[0,1],'r--')
plt.title('ROC curve: Random Forest')
plt.legend(loc='lower right')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



4.6 Top 25 words

In [41]:

```
words = tf_idf_vect.get_feature_names()
likelihood_df = pd.DataFrame(clf.feature_importances_.transpose(), columns=[ 'Score'], index=
top_25 = likelihood_df.sort_values(by='Score', ascending=False).iloc[:25]
top_25.reset_index(inplace=True)
top_words = top_25['index']
print(top_words)
```

```
0      great
1       love
2       best
3  disappoint
4     delici
5       good
6       bad
7     would
8    perfect
9       tast
10    product
11      find
12     money
13   favorit
14      like
15      use
16   thought
17     make
18    return
19   terribl
20     excel
21     nice
22    didnt
23   wonder
24   horribl
Name: index, dtype: object
```

In [42]:

```
from wordcloud import WordCloud

list_of_words_str = ' '.join(top_words)

wc = WordCloud(background_color="white", max_words=len(top_words),
               width=900, height=600, collocations=False)
wc.generate(list_of_words_str)
print("\n\nWord Cloud for Important features")
plt.imshow(wc, interpolation='bilinear')
plt.axis("off")
plt.show()
```

Word Cloud for Important features



5.Applying GBDT Algorithm

In [43]:

```

from sklearn.ensemble import GradientBoostingClassifier

Learning_rate = [0.05,0.1,0.2,0.3]
depth=[3,5,7,9,11]
param_grid = {'max_depth':depth, 'learning_rate':Learning_rate}
gb = GradientBoostingClassifier(loss='deviance',max_features='sqrt',subsample=0.1,n_estimators=100)
model = GridSearchCV(gb, param_grid, scoring = 'f1', cv=3 , n_jobs = -1,pre_dispatch=2)
model.fit(standardized_data_train,y_train)
print("Model with best parameters :\n",model.best_estimator_)
print("Accuracy of the model : ",model.score(standardized_data_test, y_test))
a = model.best_params_
optimal_learningrate = a.get('learning_rate')
optimal_depth = a.get('max_depth')

```

Model with best parameters :

```

GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.1, loss='deviance', max_depth=11,
                           max_features='sqrt', max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=100,
                           n_iter_no_change=None, presort='auto', random_state=None,
                           subsample=0.1, tol=0.0001, validation_fraction=0.1,
                           verbose=0, warm_start=False)

```

Accuracy of the model : 0.9404039665716443

In [44]:

```

results = model.cv_results_
meanscore=results['mean_test_score']

```


Heatmap for plotting CV Scores

In [45]:

```

pvt =pd.pivot_table(pd.DataFrame(model.cv_results_),values='mean_test_score', index='param_
import seaborn as sns
ax = sns.heatmap(pvt,annot=True,fmt="f")

```



In [46]:

```
clf = GradientBoostingClassifier(max_depth=optimal_depth,learning_rate=optimal_learningrate)
clf.fit(standardized_data_train,y_train)
y_pred = clf.predict(standardized_data_test)
```

5.1 Confusion Matrix

In [47]:

```
#finding out true negative , false positive , false negative and true positive
tn, fp, fn, tp = cm_tfidf.ravel()
( tp, fp, fn, tp)
print(" true negitives are {} \n false positives are {} \n false negatives are {} \n true pos
```

```
true negitives are 1139
false positives are 4975
false negatives are 67
true positives are 38819
```

5.2 Calculating Accuracy,Error on test data,Precision,Recall,Classification Report

In [48]:

```

from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import classification_report

# evaluating accuracy
acc_tfidf = accuracy_score(y_test, y_pred) * 100
print('\nThe Test Accuracy of the GBDT for maxdepth={}, Learning_rate ={} is accuracy {}'.f

# Error on test data
test_error_tfidf = 100-acc_tfidf
print("\nTest Error of the GBDT for maxdepth and Learning_rate %f%%" % (test_error_tfidf))

# evaluating precision
precision_score = precision_score(y_test, y_pred)
print('\nThe Test Precision of the GBDT for maxdepth={}, Learning_rate ={} is accuracy {}'.f

# evaluating recall
recall_score = recall_score(y_test, y_pred)
print('\nThe Test Recall of the GBDT for maxdepth={}, Learning_rate ={} is accuracy {}'.for

# evaluating Classification report
classification_report = classification_report(y_test, y_pred)
print('\nThe Test classification report of the GBDT for maxdepth and Learning_rate is \n\n

```

The Test Accuracy of the GBDT for maxdepth=11, Learning_rate =0.1 is accuracy 88.95555555555555

Test Error of the GBDT for maxdepth and Learning_rate 11.044444%

The Test Precision of the GBDT for maxdepth=11, Learning_rate =0.1 is accuracy 0.8913324410393686

The Test Recall of the GBDT for maxdepth=11, Learning_rate =0.1 is accuracy 0.9932880728282673


The Test classification report of the GBDT for maxdepth and Learning_rate is

	precision	recall	f1-score	support
0	0.84	0.23	0.36	6114
1	0.89	0.99	0.94	38886
micro avg	0.89	0.89	0.89	45000
macro avg	0.87	0.61	0.65	45000
weighted avg	0.88	0.89	0.86	45000

5.3 Plotting roc_auc curve

In [49]:

```
y_pred_proba = clf.predict_proba(standardized_data_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
plt.plot(fpr,tpr,label="tfidf, AUC="+str(auc))
plt.plot([0,1],[0,1],'r--')
plt.title('ROC curve: GBDT')
plt.legend(loc='lower right')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



5.4 Top 25 words

In [50]:

```
words = tf_idf_vect.get_feature_names()
likelihood_df = pd.DataFrame(clf.feature_importances_.transpose(), columns=[ 'Score'], index=
top_25 = likelihood_df.sort_values(by='Score', ascending=False).iloc[:25]
top_25.reset_index(inplace=True)
top_words = top_25['index']
print(top_words)
```

```
0    disappoint
1         great
2         bad
3    horribl
4        return
5         best
6        threw
7         love
8    thought
9     didnt
10    money
11    would
12    wast
13     aw
14    terribl
15    stale
16    refund
17    unfortun
18    worst
19    delici
20    good
21    pictur
22    throw
23    receiv
24    mayb
Name: index, dtype: object
```


7. Avg Word2Vec

In [53]:

```
# compute average word2vec for each review for X_train .
train_vectors = [];
for sent in sent_of_train:
    sent_vec = np.zeros(50)
    cnt_words = 0;
    for word in sent: #
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    train_vectors.append(sent_vec)

# compute average word2vec for each review for X_test .
test_vectors = [];
for sent in sent_of_test:
    sent_vec = np.zeros(50)
    cnt_words = 0;
    for word in sent: #
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    test_vectors.append(sent_vec)
```

7.1 Standardizing Data

In [54]:

```
# Data-preprocessing: Normalizing Data

from sklearn import preprocessing
standardized_data_train = preprocessing.normalize(train_vectors)
print(standardized_data_train.shape)
standardized_data_test = preprocessing.normalize(test_vectors)
print(standardized_data_test.shape)
```

```
(105000, 50)
(45000, 50)
```

7.2 Applying RandomForest Algorithm

In [55]:

```
# Gridsearch Cross Validation
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn import tree

base_learners = [40,60,80,100,120]
depth=[3,5,7,9,11]
param_grid = {'n_estimators': base_learners, 'max_depth':depth}
rf = RandomForestClassifier(min_samples_leaf=5,max_features='sqrt',criterion='gini',random_
model = GridSearchCV(rf, param_grid,scoring = 'f1',cv=3 , n_jobs = -1,pre_dispatch=2)
model.fit(standardized_data_train,y_train)
print(model.best_score_, model.best_params_)
print("Model with best parameters :\n",model.best_estimator_)
print("Accuracy of the model : ",model.score(standardized_data_test, y_test))

a = model.best_params_
optimal_estimator = a.get('n_estimators')
optimal_depth = a.get('max_depth')
```

```
0.9235738058587347 {'max_depth': 11, 'n_estimators': 120}
```

```
Model with best parameters :
```

```
RandomForestClassifier(bootstrap=True, class_weight='balanced',
                        criterion='gini', max_depth=11, max_features='sqrt',
                        max_leaf_nodes=None, min_impurity_decrease=0.0,
                        min_impurity_split=None, min_samples_leaf=5,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=120, n_jobs=None, oob_score=False,
                        random_state=100, verbose=0, warm_start=False)
```

```
Accuracy of the model : 0.9181707187599533
```


In [56]:

```
results = model.cv_results_
meanscore=results['mean_test_score']
```

Heatmap for plotting CV Scores

In [57]:

```
pvt = pd.pivot_table(pd.DataFrame(model.cv_results_), values='mean_test_score', index='param_
import seaborn as sns
ax = sns.heatmap(pvt, annot=True, fmt="f")
```



In [58]:


```
clf = RandomForestClassifier(n_estimators=optimal_estimator, class_weight='balanced')
clf.fit(standardized_data_train, y_train)
y_pred = clf.predict(standardized_data_test)
```

7.3 Confusion Matrix

In [59]:

```
cm_avgw2v=confusion_matrix(y_test,y_pred)
print("Confusion Matrix:")
sns.heatmap(cm_avgw2v, annot=True, fmt='d')
plt.show()
```

Confusion Matrix:



In [60]:

```
#finding out true negative , false positive , false negative and true positive  
tn, fp, fn, tp = cm_avgw2v.ravel()  
( tp, fp, fn, tp)  
print(" true negitves are {} \n false positives are {} \n false negatives are {} \n true pos
```

```
true negitves are 1594  
false positives are 4520  
false negatives are 416  
true positives are 38470
```

7.4 Calculating Accuracy, Error on test data, Precision, Recall, Classification Report

In [61]:

```

from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import classification_report

# evaluating accuracy
acc_avgw2v = accuracy_score(y_test, y_pred) * 100
print('\nThe Test Accuracy of the RandomForest for base_learners ={}, depth ={} is accuracy {}'.format(base_learners, depth, acc_avgw2v))

# Error on test data
test_error_avgw2v = 100-acc_avgw2v
print("\nTest Error of the RandomForest for optimal_estimator,optimal_depth is %f%%" % (test_error_avgw2v))

# evaluating precision
precision_score = precision_score(y_test, y_pred)
print('\nThe Test Precision of the RandomForest for base_learners ={}, depth ={} is accuracy {}'.format(base_learners, depth, precision_score))

# evaluating recall
recall_score = recall_score(y_test, y_pred)
print('\nThe Test Recall of the RandomForest for base_learners ={}, depth ={} is accuracy {}'.format(base_learners, depth, recall_score))

# evaluating Classification report
classification_report = classification_report(y_test, y_pred)
print('\nThe Test classification report of the RandomForest for optimal_estimator,optimal_depth is {}'.format(classification_report))

```

The Test Accuracy of the RandomForest for base_learners =120, depth =11 is accuracy 89.03111111111112

Test Error of the RandomForest for optimal_estimator,optimal_depth is 10.968889%

The Test Precision of the RandomForest for base_learners =120, depth =11 is accuracy 0.8948592695975809

The Test Recall of the RandomForest for base_learners =120, depth =11 is accuracy 0.989302062438924


The Test classification report of the RandomForest for optimal_estimator,optimal_depth is

	precision	recall	f1-score	support
0	0.79	0.26	0.39	6114
1	0.89	0.99	0.94	38886
micro avg	0.89	0.89	0.89	45000
macro avg	0.84	0.63	0.67	45000
weighted avg	0.88	0.89	0.87	45000

7.5 Plotting roc_auc curve

In [62]:

```
y_pred_proba = clf.predict_proba(standardized_data_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
plt.plot(fpr,tpr,label="Avgword2vec, AUC="+str(auc))
plt.plot([0,1],[0,1],'r--')
plt.title('ROC curve: Random Forest')
plt.legend(loc='lower right')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



8.Applying GBDT Algorithm

In [63]:

```

from sklearn.ensemble import GradientBoostingClassifier

Learning_rate = [0.05,0.1,0.2,0.3]
depth=[3,5,7,9,11]
param_grid = {'max_depth':depth, 'learning_rate':Learning_rate}
gb = GradientBoostingClassifier(loss='deviance',max_features='sqrt',subsample=0.1,n_estimators=100)
model = GridSearchCV(gb, param_grid, scoring = 'f1', cv=3 , n_jobs = -1,pre_dispatch=2)
model.fit(standardized_data_train,y_train)
print("Model with best parameters :\n",model.best_estimator_)
print("Accuracy of the model : ",model.score(standardized_data_test, y_test))
a = model.best_params_
optimal_learningrate = a.get('learning_rate')
optimal_depth = a.get('max_depth')

```

Model with best parameters :

```

GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.05, loss='deviance', max_depth=7,
                           max_features='sqrt', max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=100,
                           n_iter_no_change=None, presort='auto', random_state=None,
                           subsample=0.1, tol=0.0001, validation_fraction=0.1,
                           verbose=0, warm_start=False)

```

Accuracy of the model : 0.9425643440439163

In [64]:

```

results = model.cv_results_
meanscore=results['mean_test_score']

```


Heatmap for plotting CV Scores

In [65]:

```

pvt =pd.pivot_table(pd.DataFrame(model.cv_results_),values='mean_test_score', index='param_
import seaborn as sns
ax = sns.heatmap(pvt,annot=True,fmt="f")

```



In [66]:


```
clf = GradientBoostingClassifier(max_depth=optimal_depth,learning_rate=optimal_learningrate)
clf.fit(standardized_data_train,y_train)
y_pred = clf.predict(standardized_data_test)
```

8.1 Confusion Matrix

In [67]:

```
cm_avgw2v=confusion_matrix(y_test,y_pred)
print("Confusion Matrix:")
sns.heatmap(cm_avgw2v, annot=True, fmt='d')
plt.show()
```

Confusion Matrix:



In [68]:

```
#finding out true negative , false positive , false negative and true positive
tn, fp, fn, tp = cm_avgw2v.ravel()
(tn, fp, fn, tp)
print(" true negitves are {} \n false positives are {} \n false negatives are {} \n true pos
```

```
true negitves are 2333
false positives are 3781
false negatives are 833
true positives are 38053
```

8.2 Calculating Accuracy,Error on test data,Precision,Recall,Classification Report

In [69]:

```

from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import classification_report

# evaluating accuracy
acc_avgw2v = accuracy_score(y_test, y_pred) * 100
print('\nThe Test Accuracy of the GBDT for maxdepth={}, Learning_rate ={} is accuracy {}'.f

# Error on test data
test_error_avgw2v = 100-acc_avgw2v
print("\nTest Error of the GBDT for maxdepth and Learning_rate is  %f%%" % (test_error_avgw

# evaluating precision
precision_score = precision_score(y_test, y_pred)
print('\nThe Test Precision of the GBDT for maxdepth={}, Learning_rate ={} is accuracy {}'.f

# evaluating recall
recall_score = recall_score(y_test, y_pred)
print('\nThe Test Recall of the GBDT for maxdepth={}, Learning_rate ={} is accuracy {}'.for

# evaluating Classification report
classification_report = classification_report(y_test, y_pred)
print('\nThe Test classification report of the GBDT for maxdepth and Learning_rate \n\n ',(

```

The Test Accuracy of the GBDT for maxdepth=7, Learning_rate =0.05 is accuracy 89.74666666666667

Test Error of the GBDT for maxdepth and Learning_rate is 10.253333%

The Test Precision of the GBDT for maxdepth=7, Learning_rate =0.05 is accuracy 0.9096189702156141

The Test Recall of the GBDT for maxdepth=7, Learning_rate =0.05 is accuracy 0.9785784086817878

The Test classification report of the GBDT for maxdepth and Learning_rate

	precision	recall	f1-score	support
0	0.74	0.38	0.50	6114
1	0.91	0.98	0.94	38886
micro avg	0.90	0.90	0.90	45000
macro avg	0.82	0.68	0.72	45000
weighted avg	0.89	0.90	0.88	45000


8.3 Plotting roc_auc curve

In [70]:

```

y_pred_proba = clf.predict_proba(standardized_data_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
plt.plot(fpr,tpr,label="Avgword2vec, AUC="+str(auc))
plt.plot([0,1],[0,1],'r--')
plt.title('ROC curve: GBDT')
plt.legend(loc='lower right')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

```



9. TFIDF-Word2Vec

In [71]:

```

#tf-idf weighted w2v

from sklearn.feature_extraction.text import TfidfVectorizer

tfidf2v_vect = TfidfVectorizer()
final_counts_tfidf2v_train= tfidf2v_vect.fit_transform(x_train)
print(type(final_counts_tfidf2v_train))
print(final_counts_tfidf2v_train.shape)

final_counts_tfidf2v_test= tfidf2v_vect.transform(x_test)
print(type(final_counts_tfidf2v_test))
print(final_counts_tfidf2v_test.shape)

```

```

<class 'scipy.sparse.csr.csr_matrix'>
(105000, 38300)
<class 'scipy.sparse.csr.csr_matrix'>
(45000, 38300)

```

In [72]:

```
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(tfidf_w2v_vect.get_feature_names(), list(tfidf_w2v_vect.idf_)))

# TF-IDF weighted Word2Vec
tfidf_feat = tfidf_w2v_vect.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in sent_of_train: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            #
            tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors.append(sent_vec)
    row += 1

#Test case

tfidf_sent_vectors1 = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in sent_of_test: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            #
            tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors1.append(sent_vec)
    row += 1
print(len(tfidf_sent_vectors))
print(len(tfidf_sent_vectors1))
```

105000

45000

9.1 Normalizing Data

In [73]:

```
# Data-preprocessing: Normalizing Data
from sklearn import preprocessing
standardized_data_train = preprocessing.normalize(tfidf_sent_vectors)
print(standardized_data_train.shape)
standardized_data_test = preprocessing.normalize(tfidf_sent_vectors1)
print(standardized_data_test.shape)
```

(105000, 50)

(45000, 50)

9.2 Applying RandomForest Algorithm

In [74]:

```
#Gridsearch Cross Validation
```

```
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn import tree

base_learners = [40,60,80,100,120]
depth=[3,5,7,9,11]
param_grid = {'n_estimators': base_learners, 'max_depth':depth}
rf = RandomForestClassifier(min_samples_leaf=5,max_features='sqrt',criterion='gini',random_
model = GridSearchCV(rf, param_grid,scoring = 'f1',cv=3 , n_jobs = -1,pre_dispatch=2)
model.fit(standardized_data_train,y_train)
print(model.best_score_, model.best_params_)
print("Model with best parameters : \n",model.best_estimator_)
print("Accuracy of the model : ",model.score(standardized_data_test, y_test))

a = model.best_params_
optimal_estimator = a.get('n_estimators')
optimal_depth = a.get('max_depth')
```

0.9196179232555778 {'max_depth': 11, 'n_estimators': 120}

Model with best parameters :

```
RandomForestClassifier(bootstrap=True, class_weight='balanced',
                        criterion='gini', max_depth=11, max_features='sqrt',
                        max_leaf_nodes=None, min_impurity_decrease=0.0,
                        min_impurity_split=None, min_samples_leaf=5,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=120, n_jobs=None, oob_score=False,
                        random_state=100, verbose=0, warm_start=False)
```

Accuracy of the model : 0.9149292947964668


In [75]:

```
results = model.cv_results_
meanscore=results['mean_test_score']
```

Heatmap for Plotting CV Scores

In [76]:

```
pvt = pd.pivot_table(pd.DataFrame(model.cv_results_), values='mean_test_score', index='param_
import seaborn as sns
ax = sns.heatmap(pvt, annot=True, fmt="f")
```



In [77]:


```
clf = RandomForestClassifier(n_estimators=optimal_estimator, class_weight='balanced')
clf.fit(standardized_data_train, y_train)
y_pred = clf.predict(standardized_data_test)
```

9.3 Confusion Matrix

In [78]:

```
cm_tfidf2v = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
sns.heatmap(cm_tfidf2v, annot=True, fmt='d')
plt.show()
```

Confusion Matrix:



In [79]:

```
#finding out true negative , false positive , false negative and true positive
tn, fp, fn, tp = cm_tfdfw2v.ravel()
( tp, fp, fn, tp)
print(" true negitves are {} \n false positives are {} \n false negatives are {} \n true pos
```

```
true negitves are 1275
false positives are 4839
false negatives are 325
true positives are 38561
```

9.4 Calculating Accuracy, Error on test data, Precision, Recall, Classification Report

In [80]:

```

from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import classification_report

# evaluating accuracy
acc_tfidf2v = accuracy_score(y_test, y_pred) * 100
print('\nThe Test Accuracy of the RandomForest for base_learners ={}, depth ={} is accuracy

# Error on test data
test_error_tfidf2v = 100-acc_tfidf2v
print("\nTest Error of the RandomForest for maxdepth is  %f%" % (test_error_tfidf2v))

# evaluating precision
precision_score = precision_score(y_test, y_pred)
print('\nThe Test Precision of the RandomForest for base_learners ={}, depth ={} is accuracy

# evaluating recall
recall_score = recall_score(y_test, y_pred)
print('\nThe Test Recall of the RandomForest for base_learners ={}, depth ={} is accuracy {

# evaluating Classification report
classification_report = classification_report(y_test, y_pred)
print('\nThe Test classification report of the RandomForest for maxdepth is \n\n ',(classif

```

The Test Accuracy of the RandomForest for base_learners =120, depth =11 is a ccuracy 88.52444444444444

Test Error of the RandomForest for maxdepth is 11.475556%

The Test Precision of the RandomForest for base_learners =120, depth =11 is accuracy 0.8885023041474654

The Test Recall of the RandomForest for base_learners =120, depth =11 is accuracy 0.9916422362804094


The Test classification report of the RandomForest for maxdepth is

	precision	recall	f1-score	support
0	0.80	0.21	0.33	6114
1	0.89	0.99	0.94	38886
micro avg	0.89	0.89	0.89	45000
macro avg	0.84	0.60	0.63	45000
weighted avg	0.88	0.89	0.85	45000

9.5 Plotting roc_auc curve

In [81]:

```
y_pred_proba = clf.predict_proba(standardized_data_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
plt.plot(fpr,tpr,label="Tfidfword2vec, AUC="+str(auc))
plt.plot([0,1],[0,1],'r--')
plt.title('ROC curve: Random Forest')
plt.legend(loc='lower right')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



10.Applying GBDT Algorithm

In [82]:

```

from sklearn.ensemble import GradientBoostingClassifier

Learning_rate = [0.05,0.1,0.2,0.3]
depth=[3,5,7,9,11]
param_grid = {'max_depth':depth, 'learning_rate':Learning_rate}
gb = GradientBoostingClassifier(loss='deviance',max_features='sqrt',subsample=0.1,n_estimators=100)
model = GridSearchCV(gb, param_grid, scoring = 'f1', cv=3 , n_jobs = -1,pre_dispatch=2)
model.fit(standardized_data_train,y_train)
print("Model with best parameters :\n",model.best_estimator_)
print("Accuracy of the model : ",model.score(standardized_data_test, y_test))
a = model.best_params_
optimal_learningrate = a.get('learning_rate')
optimal_depth = a.get('max_depth')

```

Model with best parameters :

```

GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.05, loss='deviance', max_depth=9,
                           max_features='sqrt', max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=100,
                           n_iter_no_change=None, presort='auto', random_state=None,
                           subsample=0.1, tol=0.0001, validation_fraction=0.1,
                           verbose=0, warm_start=False)

```

Accuracy of the model : 0.9384902258610256

In [83]:

```

results = model.cv_results_
meanscore=results['mean_test_score']

```


Heatmap for plotting CV Scores

In [84]:

```

pvt =pd.pivot_table(pd.DataFrame(model.cv_results_),values='mean_test_score', index='param_
import seaborn as sns
ax = sns.heatmap(pvt,annot=True,fmt="f")

```



In [85]:


```
clf = GradientBoostingClassifier(max_depth=optimal_depth,learning_rate=optimal_learningrate)
clf.fit(standardized_data_train,y_train)
y_pred = clf.predict(standardized_data_test)
```

10.1 Confusion Matrix

In [86]:

```
cm_tfidf2v=confusion_matrix(y_test,y_pred)
print("Confusion Matrix:")
sns.heatmap(cm_tfidf2v, annot=True, fmt='d')
plt.show()
```

Confusion Matrix:



In [87]:

```
#finding out true negative , false positive , false negative and true positive
tn, fp, fn, tp = cm_tfidf2v.ravel()
( tp, fp, fn, tp)
print(" true negitves are {} \n false positives are {} \n false negatives are {} \n true pos
```

```
true negitves are 2040
false positives are 4074
false negatives are 887
true positives are 37999
```

10.2 Calculating Accuracy,Error on test data,Precision,Recall,Classification Report

In [88]:

```

from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import classification_report

# evaluating accuracy
acc_tfidf2v = accuracy_score(y_test, y_pred) * 100
print('\nThe Test Accuracy of the GBDT for maxdepth={}, Learning_rate ={} is accuracy {}'.f

# Error on test data
test_error_tfidf2v = 100-acc_tfidf2v
print("\nTest Error of the GBDT for maxdepth and Learning_rate  %f%%" % (test_error_tfidf2

# evaluating precision
precision_score = precision_score(y_test, y_pred)
print('\nThe Test Precision of the GBDT for maxdepth={}, Learning_rate ={} is accuracy {}'.f

# evaluating recall
recall_score = recall_score(y_test, y_pred)
print('\nThe Test Recall of the GBDT for maxdepth={}, Learning_rate ={} is accuracy {}'.for

# evaluating Classification report
classification_report = classification_report(y_test, y_pred)
print('\nThe Test classification report of the GBDT for maxdepth and Learning_rate \n\n ',(

```

The Test Accuracy of the GBDT for maxdepth=9, Learning_rate =0.05 is accuracy 88.97555555555556

Test Error of the GBDT for maxdepth and Learning_rate 11.024444%

The Test Precision of the GBDT for maxdepth=9, Learning_rate =0.05 is accuracy 0.903168302711953

The Test Recall of the GBDT for maxdepth=9, Learning_rate =0.05 is accuracy 0.9771897340945327

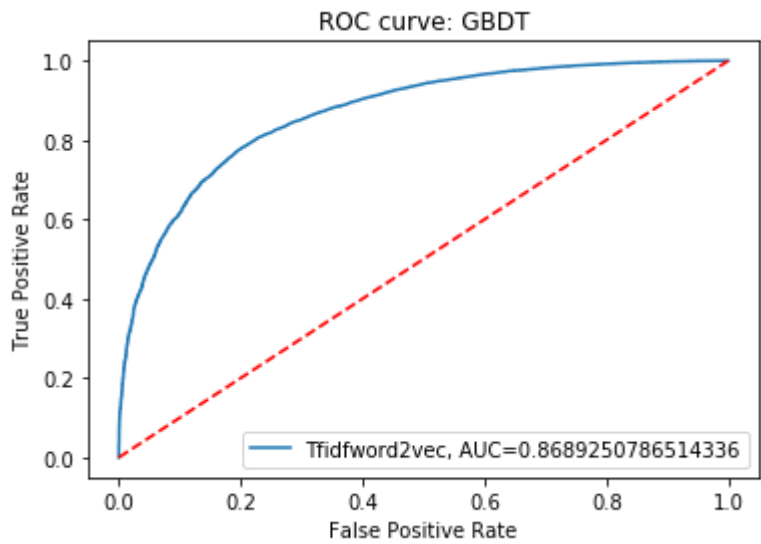
The Test classification report of the GBDT for maxdepth and Learning_rate

	precision	recall	f1-score	support
0	0.70	0.33	0.45	6114
1	0.90	0.98	0.94	38886
micro avg	0.89	0.89	0.89	45000
macro avg	0.80	0.66	0.69	45000
weighted avg	0.88	0.89	0.87	45000

10.3 Plotting roc_auc curve

In [89]:

```
y_pred_proba = clf.predict_proba(standardized_data_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
plt.plot(fpr,tpr,label="Tfidfword2vec, AUC="+str(auc))
plt.plot([0,1],[0,1],'r--')
plt.title('ROC curve: GBDT')
plt.legend(loc='lower right')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



11. Conclusion

Model Performance Table			
Model	Depth	base learner	Test Error
Accuracy			
Random Forest with Bow	11	120	11.353333
Random Forest with Tfidf	11	120	11.204444
Random Forest with Avgw2v	11	120	10.968889
Random Forest with Tfidfw2v	11	120	11.475556

Model Performance Table

Model	Depth	learning rate	Test Error
Accuracy			
GBDT with Bow	11	0.1	11.102222
88.897777			
GBDT with Tfidf	11	0.1	11.044444
88.955555			
GBDT with Avgw2v	7	0.05	10.253333
89.746666			
GBDT with Tfidf2v	3	0.05	11.024444
88.975555			

Steps Involved:-

- 1) Connecting SQL file
- 2) Data Preprocessing(Already i had done preprocessing no need to do again)
- 3) Sorting the data based on time
- 4) Taking 1st 150K Rows (Due to low Ram)
- 5) Splitting data into train and test based on time (70:30)
- 6) Techniques For Vectorization Bow,TF-IDF,word2vec,Avgword2vec,tfidfword2vec.
- 7) Normalizing Data
- 8) Applying Random Forest Algorithm
- 9) Introduced heatmap for cv_results vs max_deth vs base learner
- 10) I calculated Accuracy,Error on Test Data, Confusion Matrix,Precision Score,Recall Score,Classification Report,ROC_curve
- 11) Calculated top features and builded a word cloud
- 12) Applying GradientBoostingClassifier
- 13) Introduced heatmap for cv_results vs max_deth vs learning rate
- 12) I calculated Accuracy,Error on Test Data, Confusion Matrix,Precision Score,Recall Score,Classification Report,ROC_curve
- 13) Conclusion

In []:

